

*Electronics and Computer Science Faculty of Engineering
and Physical Sciences University of Southampton*

Josh Pattman

29 April 2023

SwarmAvg: A Novel Approach to Fully Distributed Machine Learning

Project Supervisor:

Dr Mohammad Soorati - *m.soorati@soton.ac.uk*

Second Examiner:

Dr Jo Grundy - *j.grundy@soton.ac.uk*

A project progress report submitted for the award of
Computer Science with Artificial Intelligence

Abstract

Federated learning is a technique that allows a machine learning model to be trained on data distributed across multiple data islands. This approach protects privacy by keeping the data decentralized, meaning that sensitive data does not need to ever be shared. Swarm learning is a similar technique that eliminates the need for a central server, ensuring that not only the data, but also the communication, is completely decentralized. Current Swarm Learning algorithms rely on blockchain to distribute the shared global model, however this may be a poor choice for certain scenarios closely associated with swarms. In this paper, a novel swarm learning technique called SwarmAvg is presented which operates without a blockchain. The algorithm is validated against federated learning in various scenarios, and also in some situations where federated learning cannot be applied. The benefits and drawbacks of operating Swarm Learning without a blockchain are also discussed, exemplifying some interesting reasons why one might choose to use SwarmAvg over other distributed machine learning techniques.

Statement of Originality

- I have read and understood the [ECS Academic Integrity](#) information and the University's [Academic Integrity Guidance for Students](#).
- I am aware that failure to act in accordance with the [Regulations Governing Academic Integrity](#) may lead to the imposition of penalties which, for the most serious cases, may include termination of programme.
- I consent to the University copying and distributing any or all of my work in any form and using third parties (who may be based outside the EU/EEA) to verify whether my work contains plagiarised material, and for quality assurance purposes.

You must change the statements in the boxes if you do not agree with them.

We expect you to acknowledge all sources of information (e.g. ideas, algorithms, data) using citations. You must also put quotation marks around any sections of text that you have copied without paraphrasing. If any figures or tables have been taken or modified from another source, you must explain this in the caption and cite the original source.

I have acknowledged all sources, and identified any content taken from elsewhere.

If you have used any code (e.g. open-source code), reference designs, or similar resources that have been produced by anyone else, you must list them in the box below. In the report, you must explain what was used and how it relates to the work you have done.

I have not used any resources produced by anyone else, except for a single python file submitted in the project archive. It colours the python logger so it is easier to read.

You can consult with module teaching staff/demonstrators, but you should not show anyone else your work (this includes uploading your work to publicly-accessible repositories e.g. Github, unless expressly permitted by the module leader), or help them to do theirs. For individual assignments, we expect you to work on your own. For group assignments, we expect that you work only with your allocated group. You must get permission in writing from the module teaching staff before you seek outside assistance, e.g. a proofreading service, and declare it here.

I did all the work myself, or with my allocated group, and have not helped anyone else.

We expect that you have not fabricated, modified or distorted any data, evidence, references, experimental results, or other material used or presented in the report. You must clearly describe your experiments and how the results were obtained, and include all data, source code and/or designs (either in the report, or submitted as a separate file) so that your results could be reproduced.

The material in the report is genuine, and I have included all my data/code/designs.

We expect that you have not previously submitted any part of this work for another assessment. You must get permission in writing from the module teaching staff before re-using any of your previously submitted work for this assessment.

I have not submitted any part of this work for another assessment.

If your work involved research/studies (including surveys) on human participants, their cells or data, or on animals, you must have been granted ethical approval before the work was carried out, and any experiments must have followed these requirements. You must give details of this in the report, and list the ethical approval reference number(s) in the box below.

My work did not involve human participants, their cells or data, or animals.

Acknowledgements

I would like to thank my supervisor, Dr Mohammad Soorati, for supporting and inspiring me throughout the project, and really helping me to learn a lot more about the topic of this report.

I would also like to thank my secondary supervisor, Dr Jo Grundy, for always being fast to respond to any questions I have had, and always being willing to help out.

Finally, I would like to thank my peers James Harcourt, Daniel Say, Jack Darlison, and Siobhan Tinsley for having engaging discussions with me about the direction of my project, and always providing helpful critical feedback.

Contents

1	Introduction, Problems and Goals	7
1.1	Background	7
1.2	Problems	7
1.3	Goals	8
2	Related Works	9
2.1	Federated Learning	9
2.2	Blockchain	10
2.3	Swarm Learning	10
3	Detailed Design	12
3.1	Design Overview	12
3.2	The Algorithm	13
3.3	Blockchain-less Algorithm	14
3.4	Keeping Track of Training: The Training Counter	14
3.5	Combining Neighbouring Models	15
3.5.1	Method 1: Averaging	15
3.5.2	Method 2: Averaging With Synchronisation Rate	15
3.5.3	Ignoring Old Models: Training Counter Filtering	15
3.6	Behaviour in Sparse Networks	16
3.6.1	Method 1: Passive Convergence	16
3.6.2	Method 2: Relay	17
4	Implementation	18
4.1	Machine Learning Problem	18
4.1.1	Dataset	18
4.1.2	Model	18
4.2	Implementing Federated Learning	19
4.2.1	Development of the Federated Learning Algorithm	19
4.2.2	Evaluation of this Federated Learning Implementation	19
4.3	Implementing the Prototype	19
4.3.1	Development of the Prototype	19
4.3.2	Evaluation of the Prototype	20
4.4	Implementing the Final Algorithm	20
4.4.1	Development of the Final Algorithm	20
4.4.2	Evaluation of the Implementation of the Final Algorithm	20

5	Testing Strategy and Results	21
5.1	Strategy for Gathering of Results	21
5.2	Scenario 1: Densely Connected Network	22
5.2.1	Dense Network Results	22
5.3	Scenario 2: Sparsely Connected Network	25
5.3.1	Sparse Network Results	26
5.4	Scenario 3: Densely Connected Network with Class Restrictions . . .	28
5.4.1	Dense Network with Class Restrictions Results	29
5.5	Scenario 4: Sparsely Connected Network with Class Restrictions . . .	30
5.5.1	Sparse Network with Class Restrictions Results	31
6	Critical Evaluation	32
6.1	Comparison of SwarmAvg to Other Methods	32
6.1.1	Comparison of SwarmAvg to FedAvg	32
6.1.2	Comparison of SwarmAvg to SwarmBC	33
6.2	Evaluation of this Study	33
7	Project Management	35
7.1	Time Management	35
7.2	Changing Plans	35
7.3	Risk Assessment	35
8	Conclusion and Future Work	37
8.1	Conclusion	37
8.2	Future Work	38
A	Network Generation Algorithm	41
B	Machine Learning Model	43
C	Gantt Charts	44
C.1	Gantt - Interim	46
C.2	Gantt - Final	48
D	Verifying the Word Count	49
E	Data Archive Overview	50
F	Original Project Brief	51

Chapter 1

Introduction, Problems and Goals

Machine learning is becoming an exceedingly vital tool for our society to progress. However, many modern machine learning algorithms require large volumes of diverse data to achieve optimal performance. In the ideal world, this data would be stored in a single location close to a very powerful computer for training. Unfortunately, real-world data is often distributed among multiple nodes that are unable to share the data with each other or a central location, due to privacy regulations such as GDPR [1]. Accessing a super computer for training is also a luxury that many cannot afford. The reduction in data volume available can negatively impact the post training performance [2], and the use of a slower computer means that training may not be able to be performed as fast as needed.

1.1 Background

It is evident that traditional machine learning techniques, which rely on a single machine, have certain limitations when utilised in practical applications. However, there are algorithms that aim to overcome these limitations by distributing the learning process. Moreover, a specific subset of these algorithms goes beyond distributed learning and achieves fully decentralized machine learning, which eliminates the need for a central server.

1.2 Problems

The below problems can arise from attempting to use single-machine learning techniques in the real world:

Privacy: It is common for data to be spread across multiple locations, referred to as data islands. Traditionally, all of this data would be consolidated into a single centralized server to facilitate the process of machine learning. However, it may not be possible to do so, given the potential conflict with privacy legislation. This leaves two options to the data scientist who is looking to train a model: either train a single model per data island, likely with inferior performance, or use an algorithm that would allow the different data islands to collaborate and collectively train a model.

Consider the scenario where many different hospitals wish to train a model to detect an illness in a patient. Due to the obligation to maintain patient

confidentiality, the medical data of patients cannot be shared with any of the other hospitals. This means that, despite likely having superior performance, a model trained on all data across all hospitals is not feasible, as a consequence patients would not have the highest quality medical care possible.

Performance: In general, it has been observed that larger machine learning models coupled with more data typically lead to better performance. However, the use of conventional approaches for training these large models necessitates the requirement of a powerful computer. However, most entities such as businesses or hobbyists do not have access to such a computer. Nevertheless, they may have access to multiple, lower-power computers. For instance, during non-working hours, a company may have hundreds of computers in its offices that are not being used, thus providing a pool of unused processing power.

1.3 Goals

The primary objectives of the project are delineated below. The successful completion of these objectives will be indicative of the project's overall success.

(A) Design an Novel Algorithm to Perform Machine Learning in a Swarm:

In this project, the primary aim is to design a novel algorithm to perform learning in a fully decentralised manner. This algorithm should be designed to be use in swarms of agents.

(B) Implement the Algorithm: The algorithm should be implemented in an easy-to-understand programming language with a high focus on readability, so that future work could easily expand upon it.

(C) Test Performance in Situations Where Current Approaches Perform Well:

The algorithm should be tested in situations where a current approach can be applied. The algorithm should be able to be at least comparable to the current approaches in performance.

(D) Test Performance in Situations Where Current Approaches Perform Badly:

To show that the algorithm has a potential use case, it must be shown to perform in situations where current approaches have issues. It will also be useful to show the proposed algorithm can function when certain current approaches cannot.

Chapter 2

Related Works

Machine learning is an exceedingly vital tool in modern society. Nonetheless, as the issues which machine learning endeavours to resolve become more intricate, the possibility of utilising a single centralized intelligence diminishes. One possible remedy to this problem is the distribution of both data and computation amongst multiple nodes. Transitioning from centralized approaches to distributed approaches also offers numerous advantages. For example, it has been proven mathematically that the accuracy of a distributed system surpasses that of a centralized one [3]. The following sections will provide an overview of some of the state-of-the-art approaches used to distribute machine learning.

2.1 Federated Learning

Federated Learning (FL) [4] is a technique in machine learning that aims to train a single model using all available data across nodes without requiring any data to be shared among them.

There are a multitude of published FL frameworks [5], each with different merits and drawbacks for certain use cases. Federated Averaging (FedAvg) [6] is a commonly used yet simple framework, which splits training into iterations where three steps take place:

1. A copy of the current model is sent to each node from the central server.
2. Each node performs some training with their copy of the model and their own private data.
3. The trained models from each node are sent back to the server to aggregate into the new server model.

The server model is improved over time, beyond what could be achieved by simply training on a the data stored on a single node.

As it does not require data to be shared between nodes, FL is naturally beneficial for privacy sensitive tasks compared to conventional machine learning where the data is aggregated in a central location [7]. Additionally, as FL performs training on multiple nodes in parallel, it can make better use of available training resources in

situations where processing power is not only distributed among multiple nodes but also limited on each node, such as Internet of Things (IoT) [8].

FL has been adapted into several variations which eliminate the need for a single central server. One such variant is Multi-Center Federated Learning [9], which involves multiple central servers, each connected to different clusters of nodes. Another approach is to use leader election to select a node to function as the server [10]. This method can improve network resilience because, if the server fails, a new server is elected.

2.2 Blockchain

Blockchain technology enables a network of nodes to record transactions on a shared ledger in a fully decentralized manner [11]. In a conventional blockchain, any participating node is permitted to add a transaction to the ledger, but once a transaction is executed, it becomes immutable and irreversible. Although blockchains are primarily associated with cryptocurrency and financial transactions, any type of data, including neural network weights, can be recorded as a transaction.

When a blockchain is scaled up, it may encounter performance issues. For instance, the Bitcoin network typically processes only 7 transactions per second on average [12]. High confirmation delay and performance inefficiency are two primary issues related to this paper when scaling up a blockchain [12]. These issues could prove catastrophic in a system like a swarm of robots, where it is crucial to always have the latest data and optimize limited processing power.

2.3 Swarm Learning

Swarm Learning (SL) is a subcategory of FL which operates in a completely distributed and decentralised manner. SL enables the collaboration of nodes to learn a shared global model, however in contrast to FL, a central server is never used. SL also does not use leader election, so all nodes on the network are given the same importance.

In SL, the model on which new training is performed is known as the global model. However, unlike FL where the global model is stored in a central location, the global model in SL does not materially exist, but is instead a concept which is agreed upon by the nodes in the network. An illustration of the different algorithms is shown in Figure 2.1.

One SL algorithm, referred to by this paper as SwarmBC, uses a blockchain to store the global model [13]. In this version of SL, training is performed by repeating the following steps:

1. A node obtains a copy of the global model from the blockchain.
2. The node performs some training with their copy of the model and their own private data.

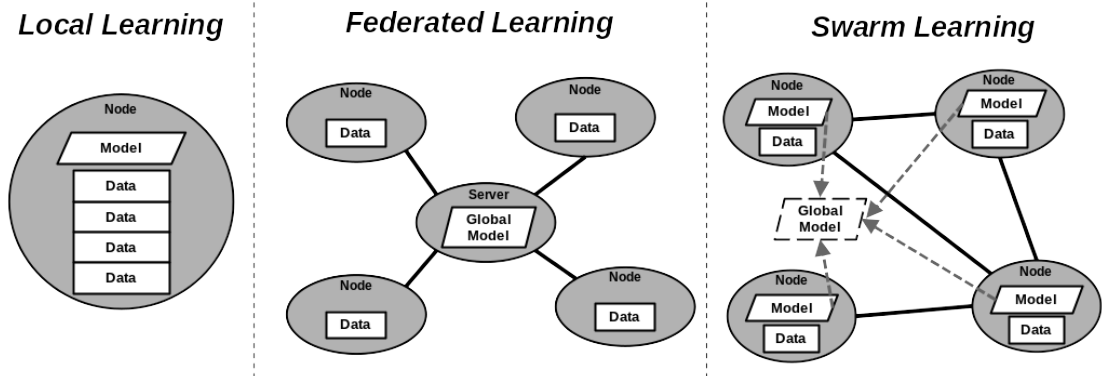


Figure 2.1: Diagram of different learning algorithms. Each *Node* indicates a single training machine, and each line denotes a connection between two machines, along which the model can be shared. In the swarm learning diagram, the dashed lines show that each local model is an approximation of the global model.

3. The updated model is merged with the latest blockchain model and sent back to the blockchain for other nodes to use.

SL exhibits many of the same benefits of FL over conventional learning, but it also improves upon FL in some aspects. As is often the case when comparing decentralised algorithms to their centralised counterparts [14], the absence of a central server theoretically makes SL more resilient and robust to failures than centralized FL approaches such as FedAvg. The absence of a central server in SL not only means that the system does not entirely rely on a single node, but also that it is less sensitive to disruptions in connections between nodes. In FL, if a connection between two nodes ceases to exist, the node can no longer participate in learning without some form of tunnelling. In contrast, in SL, even if a connection between two nodes is lost, the two connected nodes are still likely to be connected to other operational nodes, enabling the system to continue functioning normally.

The utilisation of FL with leader election effectively addresses the challenge of ensuring system robustness, as it enables any node to serve as a replacement for the server in the event of network disconnection. Nevertheless, it is important to note that a swarm of nodes is not often completely interconnected, and the connections between the nodes are often subject to changes. This dynamic and sparse network topology can present significant complexities when it comes to leader election, leading to an increase in overhead [15]. The lack of a leader election process in SL makes it a more viable option in such situations.

Finally, a swarm of nodes can offer greater scalability compared to its centralized counterpart, particularly if agents are restricted to communicating with only their nearby neighbours [16]. This is primarily because in such cases, there is no need for all communications to travel to a single node or server that may not be able to handle the significant volume of traffic. In a swarm, where nodes are limited to a certain number of close neighbours, the size of the swarm becomes less significant, since a single node would not experience a difference between a small or large swarm, as it would have the same number of neighbours. Consequently, SL is typically more scalable than FL.

Chapter 3

Detailed Design

3.1 Design Overview

In SL, a node is an agent responsible for facilitating the improvement of the global model. The global model is an abstract concept representing the consensus of all nodes in the network. In the proposed version of SL, referred to as SwarmAvg, each node maintains its own model, known as the local model, which is an approximation of the global model. However, in SwarmAvg, the consensus algorithm used is repeated averaging, not blockchain like in SwarmBC. This means that at the start of each training step, every node may start with a slightly different model. As training progresses and performance begins to plateau, each node's local model should not only converge towards a minima with respect to loss, but also towards each other.

Each node in the network possesses a confidential dataset that is not disclosed to any other nodes. In order to train the global model, nodes fit their own local model of their local dataset. In order to maintain consistency between local and global models, a combination procedure is conducted following each round of local training, which involves the integration of neighbouring nodes models into the local model.

SwarmAvg works by repeating the training step until training is complete. The actions in each training step for SwarmAvg are as follows:

1. Fit local model to local dataset.
2. Send local model to all neighbours.
3. Combine neighbouring models into local model, using one of the methods presented in Section 3.5.

Additionally, each node retains a local cache of the most recent models of its neighbouring nodes. This cache is updated each time a neighbouring node transmits its model to the node in question, instead of being updated on-demand during the combination step. The reasoning behind this decision is explained further in Section 4.3.1.

3.2 The Algorithm

The training step, found at Algorithm 1, is the section of the algorithm which runs continuously during the time which a node is running. It takes care of training and synchronising the local model. The provided code represents what happens in a single training step, meaning that it should be run in a loop that terminates once a stop condition, such as target accuracy, has been reached.

Algorithm 1 A Single Training Step - should be called repeatedly in a loop

```

1: TRAIN(localModel, localData)
2: for all  $n \in neighbors$  do
3:   SENDTO((localModel, localTrainingCounter),  $n$ )
4: end for
5: for  $x \in range(maxSyncWaits)$  do
6:    $neighborModels \leftarrow \emptyset$ 
7:   for all  $n \in neighbors$  do
8:      $model, trainingCounter \leftarrow CACHELOOKUP(n)$ 
9:     if  $trainingCounter + \beta \geq localTrainingCounter$  then
10:      APPEND( $neighborModels$ ,  $model$ )
11:    end if
12:  end for
13:  if  $length_{neighborModels} \geq \gamma$  then
14:    if  $synchronisationMethod = "AVG"$  then
15:       $localModel \leftarrow \mu(localModel \cup neighborModels)$ 
16:    else if  $synchronisationMethod = "ASR"$  then
17:       $localModel \leftarrow (1 - \alpha) * localModel + \alpha * \mu(neighborModels)$ 
18:    end if
19:  else
20:    continue
21:  end if
22:  SLEEP(syncWaitTime)
23: end for

```

The model received event, which can be found at Algorithm 2, is run on the local node every time a remote node sends the local node a model update. This event takes care of updating the local model cache, to ensure the local node has the most up-to-date information. The model update sent from the remote node should contain the model and training counter of the remote node.

Algorithm 2 Model Received Event - called when a model update is received from a remote node

Input: *neighbour, nModel, nTrainingCounter*

```
1: if INCACHE(neighbour) then
2:    $\_, nTraningCounterOld \leftarrow \text{CACHELOOKUP}(n)$ 
3:   if nTrainingCounter > nTraningCounterOld then
4:     SETCACHE(neighbour, (nModel, nTrainingCounter))
5:   end if
6: else
7:   SETCACHE(neighbour, (nModel, nTrainingCounter))
8: end if
```

3.3 Blockchain-less Algorithm

Whereas SwarmBC employs a blockchain as the mechanism for distributing the global model, SwarmAvg utilises a variation of averaging with its neighbours and does not use a blockchain. This decision was made due to the long transaction confirmation time of large blockchains, which could adversely affect training as nodes continuously upload their latest networks to the network. If this process were to take an excessive amount of time, each node would be training on an outdated model.

Furthermore, the inefficiency of blockchains with respect to performance is another reason for not using them in SwarmAvg. SwarmAvg is designed for deployment in large swarms, with each agent possibly having low processing power. If a blockchain were utilised, some of the processing power that could be devoted to training would instead be utilised for validating transactions. In scenarios where processing power is limited, it would be practical to avoid the use of blockchains.

3.4 Keeping Track of Training: The Training Counter

A vital aspect of SwarmAvg, specifically the combination step, involves evaluating the performance of a nodes local model. The conventional approach would involve testing each model using an independent test set. However, due to the inability to exchange test sets among nodes, this approach is not feasible as it would result in non-comparable scores for each model. In order to circumvent this problem, this paper presents a heuristic metric referred to as the "training counter," which serves as an approximation of the level of training for a network by estimating the number of training steps performed on a given model.

The training counter can be changed in one of two manners. Firstly, the counter is incremented by 1 when the local model is trained on the local dataset, indicating that an additional step of training has been performed. Following the combination step, the training counter is also updated to reflect the combination method that was utilised. For instance, if the neighbouring models were averaged, the training counter would be updated to represent the average of all neighbouring nodes' training counters.

3.5 Combining Neighbouring Models

The combination step is a crucial component of SwarmAvg. During this step, a node merges its local model with those of its neighbours, producing an updated estimate of the global model. This paper presents multiple methods for performing the combination step.

In the below equations, $\mu(x)$ denotes the function *mean*(x). All models are assumed to be flattened into 1-dimensional arrays of parameters.

3.5.1 Method 1: Averaging

The most rudimentary approach to combination is to compute the average model between the local model and the models of all neighbouring nodes.

$$localModel \leftarrow \mu(localModel \cup neighborModels)$$

This technique is utilised in FedAvg, which is the most simplistic form of FL. The benefit of this method is that it necessitates no hyper-parameters, which means that the data scientist needs to do less tuning. However, this attribute can also be viewed as a drawback, as it affords less flexibility in terms of customization for particular tasks.

3.5.2 Method 2: Averaging With Synchronisation Rate

A more complex approach to combination is to compute the average model of all neighbours, then compute the weighted average between that model and the local model.

$$localModel \leftarrow (1 - \alpha) * localModel + \alpha * \mu(neighborModels)$$

The synchronisation rate, denoted as α , indicates the degree to which each node adjusts its local model to align with the global model. If α is set too low, each node's model in the network will diverge, resulting in each node becoming trapped at a local minima, an effect referred to as divergent training. On the other hand, if α is too high, the progress achieved by the local node will be discarded at each averaging step, which can result in slower learning.

This method is beneficial over averaging, as described in Section 3.5.1, as it provides some resistance to variation in the number of neighbours of a node. When using averaging, the amount of change made by the local node that persists into the next training step is proportionate to the number of neighbours which were combined. However, when using averaging, the amount of change made that persists will be constant, no matter how many neighbours that are present. This may increase the stability of training, especially if the situation which SwarmAvg is deployed involves a dynamic set of neighbours.

3.5.3 Ignoring Old Models: Training Counter Filtering

A potential modification to both of the previously mentioned combination algorithms involves filtering based on the training counter. Specifically, a node may only include

its neighbour models if they meet the following statement:

$$neighborTrainingCounter + \beta \geq localTrainingCounter$$

The training offset β is the amount the training counter of a neighbour can be behind the local training counter before it is ignored.

A compelling reason to allow training counter filtering is the increased fault tolerance. Consider the situation where node A has received many model updates from many nodes, one of which being node B . However, node B goes offline and no longer is sending model updates. Without training counter filtering, node A will continue to combine the outdated node B model with it's own for as long as it is training. However, if training counter filtering is enabled, after a number of training steps the outdated model B updates will be ignored.

An issue with training counter filtering pertains to the presence of runaway nodes. These nodes possess a substantially higher training counter compared to all other nodes in the network, meaning that when filtering is applied, they are left with no neighbours to utilise in the combination step. Consequently, a runaway node may start to overfit on its own training data, as this is the only data it is exposed to, thereby leading to decreased local performance, as well as potential performance reductions in the rest of the network. To address this problem in SwarmAvg, each node must wait until it has obtained at least γ viable neighbours prior to performing the combination step. Although this measure prevents individual nodes from becoming runaway nodes, groups of size γ still have the potential to become runaway as a unit. Nevertheless, if γ is roughly equivalent to the number of neighbours and all neighbours train at a comparable rate, the issue is minimised.

3.6 Behaviour in Sparse Networks

Given the sparsely connected nature of distributed scenarios, it is often the case that nodes only have direct connections to a small subset of their neighbours. Although FL struggles in this situation, the author hypothesises that SwarmAvg should be able to function in sparsely connected networks.

3.6.1 Method 1: Passive Convergence

An approach to deal with a sparsely connected network is to use the swarm learning algorithm without any modifications. This approach is effective due to the use of averaging as a combination method. When a node tries to update the global model, its changes will propagate through the network slowly, over many training iterations, even to nodes that are not directly connected. This approach has the advantage of requiring no extra data transmission, resulting in significantly less data traffic compared to other methods.

However, this method also has certain theoretical drawbacks. Consider a scenario where the network is comprised of several sparsely connected groups of nodes, where each node in a group is densely connected to other nodes within that group. In this case, it is possible that each group may learn a distinct solution to the problem. This

is inefficient because instead of functioning as a cohesive network, there are multiple smaller networks acting somewhat independently of each other, potentially leading to a decrease in overall performance.

3.6.2 Method 2: Relay

A solution to the aforementioned divergence of distant groups could be to relay any received model updates, meaning that as long as each node has at least one path to reach all other nodes, the network will behave as a densely connected network. This approach offers theoretical immunity to changes in network topology, but in practice, the network's performance may still decrease compared to a truly dense network due to slower communication times between non-connected nodes.

The main disadvantage of this approach is the drastic increase in network traffic, which in turn will lead to longer model transfer times. If the swarm learning algorithm is applied to a low power network, such as an IoT network, the increase in network traffic may not be feasible at all. This approach can also be applied to FL with the same advantages and disadvantages. For these reasons, relaying model updates will not be discussed further.

Chapter 4

Implementation

4.1 Machine Learning Problem

In order to evaluate SwarmAvg, it was necessary to define both a problem and a model to supply the algorithm with updates. Nevertheless, it should be noted that the algorithm is intended to function with any model and dataset. Presented below are the selected model and dataset.

4.1.1 Dataset

Initially, the dataset utilised for experimentation was the MNIST dataset, which encompasses 60,000 greyscale 28x28 labelled images of digits from 0 to 9. This dataset was selected for its simplicity, requiring no pre-processing or data cleaning prior to training, and due to its availability as a built-in component of the chosen machine learning framework, Keras. It has also been used to test FL frameworks in past research [17].

Conversely, upon implementation it was determined that this dataset was too simple for the application of machine learning, as a single node could reach near peak accuracy after a single epoch, rendering it ill-suited for swarm learning, an algorithm designed to function across multiple training epochs. To address this issue, the MNIST dataset was replaced with MNIST-fashion (henceforth referred to as MNIST-F), a drop-in replacement dataset containing 10 classes of different items of clothing. MNIST-F is known to be more challenging [18].

4.1.2 Model

The Keras machine learning framework in Python was utilised to implement the model due to its reputation for being both simple and straightforward. All experiments made use of the same model, a small convolutional neural network, which is outlined in Appendix B. The model was tested on the MNIST-F dataset and was able to attain an accuracy score of above 90 percent when trained on the entire dataset using local learning; this result is on par with the accuracy reported in the original paper [18], meaning that the model was complex enough to learn the dataset, and therefore was suitable for use during testing.

4.2 Implementing Federated Learning

FedAvg was chosen for comparison of performance against SwarmAvg. In order to ensure fairness of the comparison, it was necessary to implement FedAvg from scratch using the same language framework as SwarmAvg.

4.2.1 Development of the Federated Learning Algorithm

The implemented algorithm worked in the same manner as the algorithm described in the FedAvg paper, although one modification was made: at the start of each timestep instead of choosing N random nodes to perform training, all nodes were chosen. This means that every available node will perform training at every timestep, which should result in the best possible performance for federated learning, especially given that only 10 nodes could be run at once.

Initially, a REST API was utilised to transfer the model between the server and the client. However, this was deemed unnecessary and was supplanted for two reasons. Firstly, the decision was made to measure performance against training steps instead of performance against time, which is discussed in further detail in the results section. Secondly, the REST API approach was much slower than the method chosen to replace it, yet it resulted in the same performance measurements when gauged in terms of training step. As a substitute for the REST API, a system of functions was implemented. When a node sent a model to another node, it simply called a function on that node, however this simplification was abstracted away from the main algorithm code. The faster training enabled a greater number of experiments to be conducted.

4.2.2 Evaluation of this Federated Learning Implementation

This implementation of federated learning is simple yet effective. Though certain simplifications have been made, they should not interfere with the performance of the model in this scenario. It should only be used for testing performance against training steps, not against time taken, as the model transfer layer has been simplified.

4.3 Implementing the Prototype

The decision was made to make an initial, less streamlined, prototype as a proof-of-concept before spending a large amount of time creating the final algorithm. This prototype was created to be very modifiable so that changes could easily be made and tested.

4.3.1 Development of the Prototype

This algorithm was a simplified version of that described in the design section. The primary difference was that when a node performed its synchronisation step, it would request the models from each neighbour instead of utilising its cached versions of their models. This had the consequence of slowing down training, as the synchronisation

step could not be completed until all nodes had responded. Furthermore, this algorithm did not incorporate the training counter, leading to the absence of training counter filtering, β and γ .

4.3.2 Evaluation of the Prototype

This step was beneficial for the progression of the project, as it enabled the author to form the ideas detailed in the design process, which were then implemented in the subsequent step. However, due to the abundance of superfluous code, the algorithm was inefficient and performed poorly. For this reason, it was decided not to record the results of this method.

4.4 Implementing the Final Algorithm

In this implementation, the findings of the prototype were taken and built upon. The code was streamlined for the purpose of testing performance. However, the code was still designed to be reusable and easy to read.

4.4.1 Development of the Final Algorithm

The algorithm is as described in Section 3. However, there was one discrepancy which needed to be tested: whether to send model updates to neighbours before or after performing the combination step. Both were tested, but pushing model updates before combination seemed to be the more effective method when comparing the accuracy.

The back end for distributing models was implemented as an interface which abstracts the details of the distribution away from the main algorithm code. For this implementation, the same distribution strategy as the previously implemented FedAvg was used: calling local functions that simulate a web connection.

4.4.2 Evaluation of the Implementation of the Final Algorithm

Overall, this implementation of SwarmAvg is satisfactory. It is efficient, reusable and simple to understand and use. Nevertheless, it is not yet suitable for real world applications, as it lacks any security features and error handling, meaning that cyber attacks could be performed with ease. However, it would not be challenging to incorporate a secure, internet-enabled back-end in the future.

Chapter 5

Testing Strategy and Results

5.1 Strategy for Gathering of Results

The experiments were conducted through the simulation of a group of nodes on a single computer. The accuracy of each node was evaluated on an unseen test set and subsequently recorded in a file after each training step. All nodes shared the same test set. Below are the specific details of the running of each experiment:

Chosen Metric: Accuracy: The metric chosen for the following experiments was accuracy, due to its comprehensible nature. Despite the fact that an unbalanced dataset presents one of the significant drawbacks of using accuracy, this concern is irrelevant in the case of MNIST-F since it is a balanced dataset, meaning it has the same number of samples in each class. The accuracy of each node is computed after each training step using the test subset of MNIST-F, and none of the nodes are ever provided access to the test set for training.

Number of Simulated Nodes: The experiments were conducted using 10 nodes, with the exception of the server in cases where FedAvg was employed. The decision of how many nodes to simulate was based on the highest node count attainable without causing inconsistencies and crashes due to resource depletion of the training machine.

Repeated Testing for Noise Reduction: The training process for each experiment was conducted five times, and the resulting accuracies of every node were recorded. To mitigate the impact of training noise on the performance graphs, the accuracy value for each time step was calculated as the median accuracy across all nodes and runs at that time step. As there were 10 nodes, this resulted in the graphs representing median of 50 nodes.

Configuring the Algorithm: In each of the following experiments, the algorithm was configured using a specific set of parameters $\alpha\beta\gamma$. These parameters were obtained heuristically by making an initial guess, testing, and then fine-tuning them until a satisfactory outcome was reached. However, it is important to note that an exhaustive investigation into the optimal parameter configuration for a particular type of problem is not within this paper's scope, meaning that it is plausible that swarm learning could yield better results with more precisely tuned parameters.

Data Provided to Each Node: The experiments evaluate the algorithm’s performance using three levels of data volume per node. These levels are considerably smaller than the full MNIST-F dataset not only to increase problem difficulty, but also as the algorithm is intended for scenarios where each nodes access to data is restricted. To create a subsection of data for each node, a random sampling with replacement method was used to select the desired number of datapoints. During the initial training phase, each node performs a single sampling of its dataset, after which that nodes data subset remains constant. The three levels of data volume can be seen in Table 5.1.

Training Steps and Epochs: Due to the limited size of the dataset, a single node executes more than one epoch of training in each training step. The number of epochs carried out by a node per training step will be referred to as Epochs per Step (EPS). Empirical testing has indicated that both SL and FL exhibit improved performance with higher EPS, at times surpassing the gains from increasing the number of training steps. Moreover, the utilisation of higher EPS was favoured due to its reduced training time, compared to increasing the number of training steps. The three levels of data volume with their respective EPS are shown in table 5.1.

Dataset Size	EPS
1000	5
100	10
25	20

Table 5.1: The different levels of dataset size and EPS that were tested


5.2 Scenario 1: Densely Connected Network

A crucial experiment for evaluating the performance of the SwarmAvg algorithm involves assessing its performance under optimal circumstances, specifically within a network of nodes wherein each node is directly connected to every other node. In FedAvg, the analogous topology involves direct connections between each node and the server. This comparison is significant as it facilitates a direct evaluation of the SwarmAvg and FedAvg algorithms under their respective ideal conditions.

5.2.1 Dense Network Results

The following are the results obtained from the execution of the training script. Each graph displays the data for SwarmAvg in red, and FedAvg in black.

Accuracy by Training Step for 1000 Samples




(a) Full Graph

(b) Zoomed Graph

Figure 5.1: Median accuracy by training step across 5 repeats. Each node has 1000 random data samples from MNIST-F. The shaded region represents the upper and lower quartile.

The results depicted in Figure 5.2 demonstrate that the SwarmAvg algorithm exhibits a slower convergence rate compared to FedAvg, trailing by at least 1 epoch for the duration of the test. Despite this, both methods ultimately achieve a similar level of accuracy, with less than a percent difference.

Accuracy by Training Step for 100 Samples



(a) Full Graph

(b) Zoomed Graph

Figure 5.2: Median accuracy by training step across 5 repeats. Each node has 100 random data samples from MNIST-F. The shaded region represents the upper and lower quartile.

Accuracy by Training Step for 25 Samples

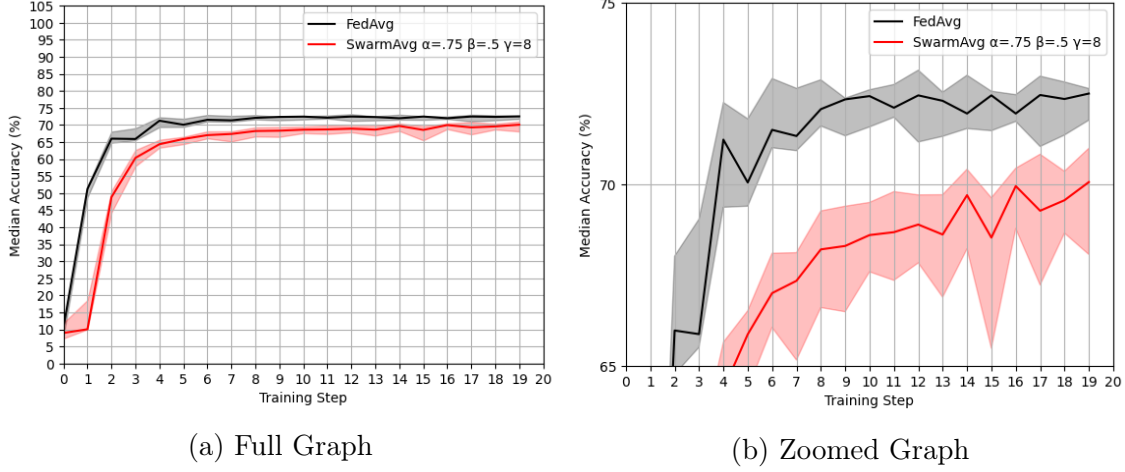


Figure 5.3: Median accuracy by training step across 5 repeats. Each node has 25 random data samples from MNIST-F. The shaded region represents the upper and lower quartile.

The trend of SwarmAvg ending with a lower accuracy than FedAvg continues in Figure 5.3. Both algorithms also have a much lower accuracy than what they achieved previously, due to the drastic decrease in available data.

Similarly to Figure 5.2, Figure 5.3 shows that SwarmAvg concludes with a lower accuracy than FedAvg. Furthermore, both algorithms exhibit a considerable decline in accuracy when compared to their previous performances with higher volumes of data.

The most noticeable impact resulting from reducing the volume of data is the significant decrease in the accuracy of both algorithms, as expected. However, it is also evident that the reduction in data affects SwarmAvg and FedAvg to a similar degree, as both of their peak accuracies stay within a similar range of each other in all tests. The difference in peak accuracy between the two algorithms was quite small, often within a 2 percent margin.

One of the prominent challenges associated with SwarmAvg is its slower convergence rate compared to FedAvg, particularly from the outset. SwarmAvg consistently takes a longer time to attain its peak accuracy. This may be attributed to the asynchronous nature of the nodes in SwarmAvg, which implies that some nodes that conduct training before others may have lower accuracy than expected.

It is worth noting that in all these evaluations, SwarmAvg has a slightly higher inter-quartile range than FedAvg, indicating that FedAvg is more consistent in terms of accuracy. However, this difference is minor.

One interesting feature of both SwarmAvg and FedAvg is that both algorithms seem to be affected by overfitting much less than would be expected. Especially with just 25 data samples per node, the effects of overfitting would usually be far more severe in a single node. However, this does not mean that the algorithms do not overfit, but it may indicate that the process of overfitting for both algorithms is drastically slowed down.

5.3 Scenario 2: Sparsely Connected Network

In reality, it is uncommon for each node to be linked with every other node. To simulate a more realistic scenario, a technique was employed to generate a network of nodes with a specific number of connections. When density is set to 0, the network is minimally linked, meaning that each node has at least one indirect path to every other node, but the minimal number of connections required to accomplish this exist. When density is set to 1, all nodes are connected to one another in a dense fashion. Since this measure may not provide a straightforward indicator of network density, two additional metrics will be provided: Mean Minimum Hops (MMH) and Mean Connections per Node (MCPN). MMH denotes the mean minimum number of transitions required to get from one node to another in the network. MCPN denotes the average number of connections a given node possess, for a network of 10 nodes.

In order to conduct thorough testing on a variety of potential deployment scenarios, several different densities were tested for each count of data. The densities that were tested, along with their corresponding MMH and MCPN, are presented in Table 5.2.

Density	MMH	MCPN
1	1.0	9.0
0.75	1.2	7.2
0.5	1.4	5.4
0.25	1.7	3.6
0	3.0	1.8

Table 5.2: The statistics for different density levels

In order to assist the reader with visualizing the various density levels, some sample networks that were generated for each density can be found in Figure 5.4.




Figure 5.4: Example networks of nodes generated for each density level, visualised using the tool at https://csacademy.com/app/graph_editor. Each time a simulation is started, a new random network is generated for that simulation.

This section does not include any testing of FedAvg. In scenarios where not every node is directly connected to the server node, FedAvg has two potential options: ignore all nodes which are not directly connected, or attempt to relay the model updates through connected nodes. As mentioned in Section 3.6.2, relaying may not always be the optimal solution. Ignoring nodes is also not a good option, as data is wasted. Therefore, in this test, the decision was made to solely evaluate SwarmAvg.

5.3.1 Sparse Network Results

Below are results obtained from testing. The density of each line on the graph is indicated by the colour, with a green hue representing higher density and a red hue indicating lower density. For all tested densities, the value of γ was set to $\text{round}_{\text{down}}(MCPN) - 1$, which ensures that each node can progress only after waiting for all but one of its neighbours. Notably, failure to reduce γ for less dense networks would cause several nodes to wait for a number of neighbours that cannot be achieved. Due to the random nature of the graph generation, there will still be some nodes who have less than γ neighbours, but this should be rare and the loop to wait for γ neighbours will terminate after a certain number of tries anyway, ensuring that training can progress.

Accuracy by Training Step for 1000 Samples, with Varying Density

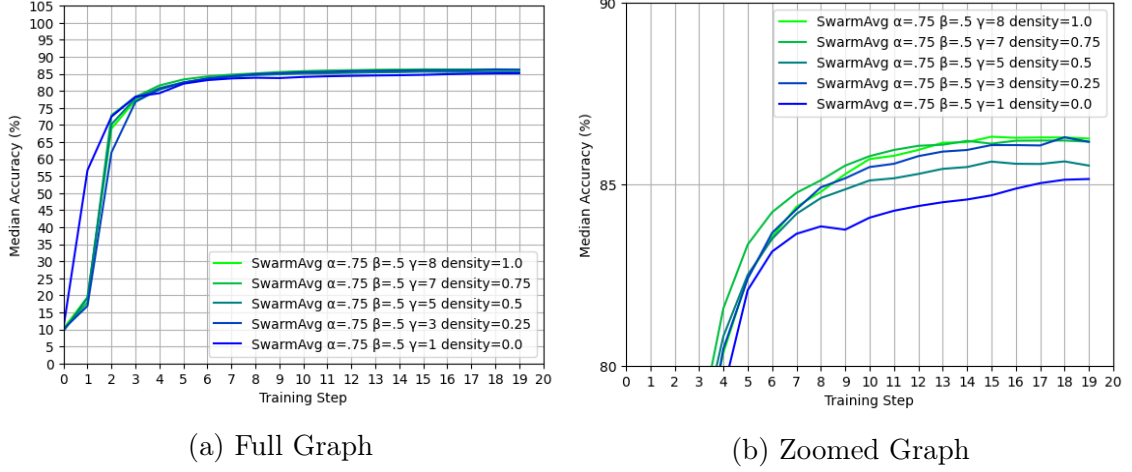


Figure 5.5: Median accuracy by training step across 5 repeats. Each node has 1000 random data samples from MNIST-F. Quartiles are not shown.

An observation that can be made of Figure 5.5 is that the different network densities perform very similarly. The final accuracy for all of the tests fell within a 2 percent range, with the higher densities occupying the higher end.

One interesting feature of the Figure 5.5 is that the lowest density demonstrates faster convergence initially than the other densities. The author hypothesises that this may be due to groups of nodes forming local sub-networks which converge on their own sub-sets of data faster, but more slowly combine into the network as a whole as training progresses.

Accuracy by Training Step for 100 Samples, with Varying Density

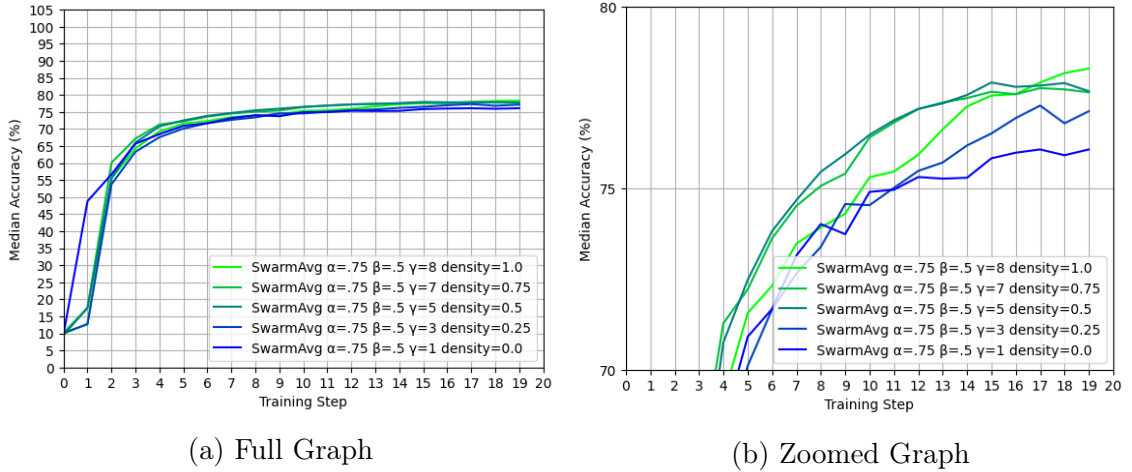


Figure 5.6: Median accuracy by training step across 5 repeats. Each node has 100 random data samples from MNIST-F. Quartiles are not shown.

In Figure 5.6, there is a greater spread of final accuracies, about 2.5 percent, than is shown in Figure 5.5. The more noticeable feature of Figure 5.6 is that the final accuracies of all densities are approximately 8 percent lower than their counterparts with more data available in figure 5.5.

Accuracy by Training Step for 25 Samples, with Varying Density

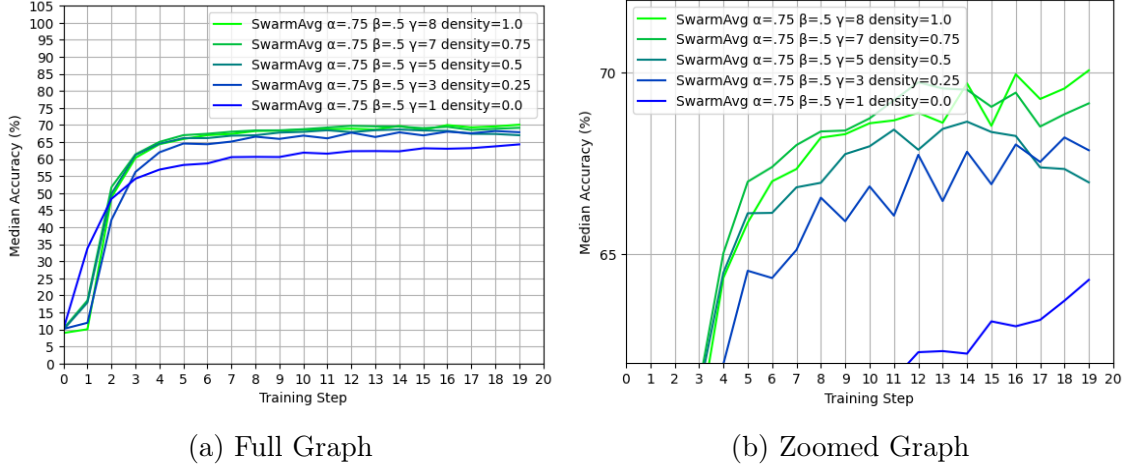


Figure 5.7: Median accuracy by training step across 5 repeats. Each node has 25 random data samples from MNIST-F. Quartiles are not shown.

With 25 data samples, the tests depicted in Figure 5.7 show a much higher variation in final accuracies than the previous two figures. There is also a much higher degree of noise in this figure than any other, suggesting that training may be much less stable with this small number of data samples to train on.

It should also be noted that the lowest density is still exhibiting the behaviour of outperforming the other densities in the initial steps, however it appears that this effect has become slightly less prominent.

Generally, altering the network density of nodes has a small impact on the training of the nodes within the network. Decreasing the density of nodes typically leads to a lower final accuracy. This effect becomes more noticeable as the data volume provided to each node decreases, as demonstrated by the increased variability in training displayed in Figure 5.7 in comparison to Figures 5.5 and 5.6.

The networks possessing a density of 0 attain their optima at a faster rate than those with higher densities. Figure 5.7 illustrates that the lower densities reach convergence marginally quicker than higher densities, yet are surpassed by the latter towards the end of training.

5.4 Scenario 3: Densely Connected Network with Class Restrictions

Decentralized machine learning algorithms, such as FedAvg and SwarmAvg, commonly involve a scenario where each node possesses a dataset whose distribution does not align with the overall data distribution across all nodes. In order to simulate such a scenario, a situation was created where each node was given access a unique set of three out of ten classes of MNIST-F. The tests in this experiment were conducted in a densely connected network of nodes, thereby allowing the evaluation of FedAvg as well.

5.4.1 Dense Network with Class Restrictions Results

Accuracy by Training Step for 1000 Samples of 3 Classes




Figure 5.8: Median accuracy by training step across 5 repeats. Each node has 1000 random data samples from MNIST-F. Quartiles are depicted as the shaded area. Each node has access to a unique set of 3 classes out of 10

During the course of the experiment, it was determined that SwarmAvg required some adjustments in order to function properly. One issue was a significant fluctuation in accuracy during the training process, which may have been due in part to the synchronization of nodes in the network. To address this concern, the decision was made to decrease the value of γ and introduce a delay in the node start times, in order to desynchronize their training. The chosen delay time was 2 seconds. Figure 5.8 illustrates these two techniques, with the delayed training and reduced gamma method depicted in red, and the original method depicted in green. Evidently, these adjustments had a positive impact in minimizing the oscillations, although they were not entirely eliminated.

While attempting to adjust the parameters, another issue was encountered was divergent training, which occurs when nodes begin to learn completely separate models and start to disregard the models of their neighbours. It was observed that a low value of α was sufficient to cause this problem. When divergent training occurred, the nodes failed to achieve an accuracy level higher than 30 percent, which is expected if a node only trains on three out of the ten available classes. This effect is shown as the blue line of Figure 5.8.

It is evident that SwarmAvg consistently exhibits a significantly lower level of accuracy compared to FedAvg. Additionally, the inter-quartile range of SwarmAvg is noticeably higher.

Accuracy by Training Step for 100 and 25 Samples, with 3 Classes

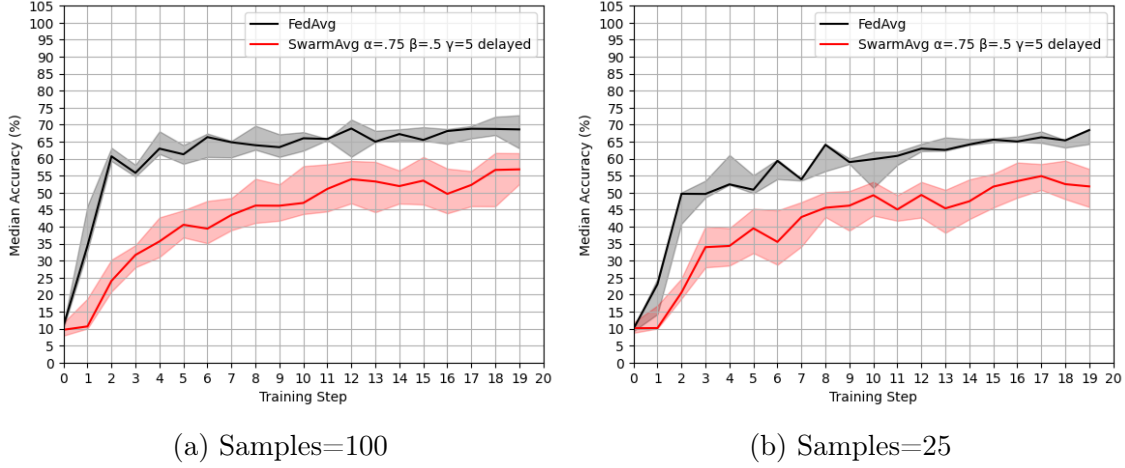


Figure 5.9: Median accuracy by training step across 5 repeats. Quartiles are depicted as the shaded area. Each node has access to a unique set of 3 classes out of 10

When performing training on a lower sample count, as depicted in Figure 5.9, the oscillations in SwarmAvg seem to have been drastically reduced when compared to 5.8. However, both SwarmAvg and FedAvg have lower accuracies by the end of training than they did with 1000 samples.

It is clear that the reduction of data volume per node has a detrimental impact on the performance of SwarmAvg in general. However, in these restricted class tests, this effect is less pronounced compared to previous tests where each node had access to all possible classes. It should be noted that SwarmAvg encounters accuracy oscillations, which can be partially attributed to the synchronized training steps of nodes. Nevertheless, by decreasing the value of γ and staggering the startup of nodes, these oscillations can be mitigated. Moreover, reducing the data volume also seems to aid in reducing the effect of these oscillations, possibly implying that nodes are being trained excessively in a single step.

5.5 Scenario 4: Sparsely Connected Network with Class Restrictions

This scenario was designed to test the SwarmAvg algorithm to its limits. It not only involved limiting data to each node, but also limiting the classes each node had access to in the same fashion as Scenario 3, and also constraining the networks to be sparsely connected.

5.5.1 Sparse Network with Class Restrictions Results

Accuracy by Training Step for 1000, 100 and 25 Samples, with 3 Classes and Varying Density

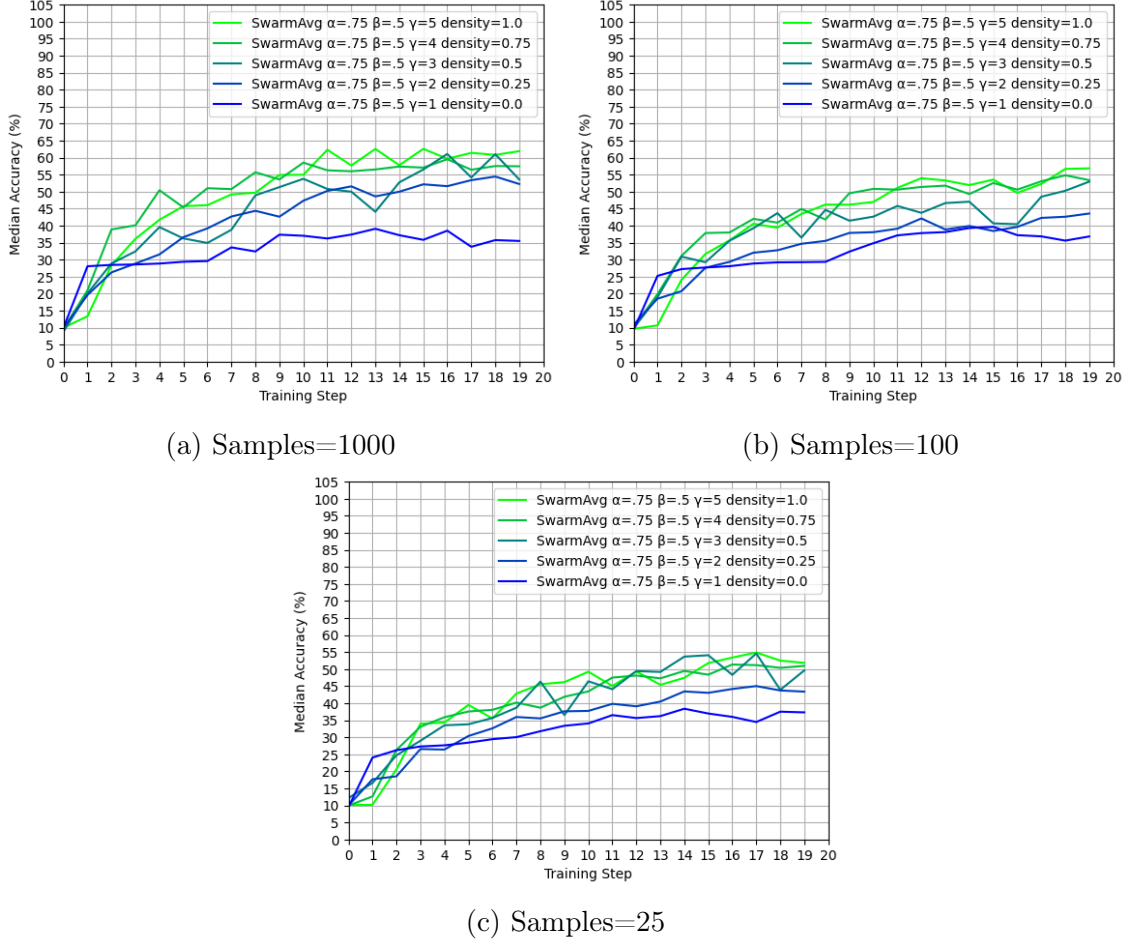


Figure 5.10: Median accuracy by training step across 5 repeats. Quartiles are not shown. Each node has access to a unique set of 3 classes out of 10.

In all three tests, it is clear that a higher density is highly correlated with higher performance in this scenario. It is also noteworthy that with a network density of 0, no matter which volume of data was presented, the accuracy curve was very similar. However, as the density of the network increased, so did the degree at which it was effected by the reduction in data. This scenario also presented graphs with a high degree of noise when compared to other scenarios, despite representing the median of 5 training runs. However, the level of noise is somewhat consistent on all three tests with different data volumes.

Chapter 6

Critical Evaluation

6.1 Comparison of SwarmAvg to Other Methods

Although SwarmAvg has demonstrated promising results, it remains important to conduct a comparative analysis against established methods. This will enable individuals seeking to incorporate distributed machine learning into their future projects to make an informed decision whether SwarmAvg is the algorithm that they should use.

6.1.1 Comparison of SwarmAvg to FedAvg

SwarmAvg has the significant benefit of being more fault tolerant than FedAvg. This is because FedAvg has a central server, which means that the whole network would stop functioning if the central server is stopped. Not only this, but if a node drops its connection to the server, that node is effectively ignored from training, which has a negative effect on performance.

However, SwarmAvg has some significant drawbacks. For example, SwarmAvg has been shown to be slower to converge than FedAvg. Not only this, but SwarmAvg can create a significantly higher amount of network traffic than FedAvg or FL in general. This effect is maximised when every node is connected to every other node, meaning that the number of connections scale by $O(n^2)$, where n is the number of nodes. FL will always both number of connections and network traffic scale with $O(n)$, as every node only needs a single connection to the server. However, it is worth noting that SwarmAvg will consume less network traffic in less dense networks of nodes.

An important advantage of SwarmAvg over FedAvg is that former can be used in sparsely connected networks without any modification. This means that it is possible for it to SwarmAvg to run even when most nodes in the network are not connected to most other nodes, which may be a more realistic scenario for certain low powered networks of devices such as IoT networks or robotic swarms. For this reason, SwarmAvg may be a better choice than FedAvg for these situations.

6.1.2 Comparison of SwarmAvg to SwarmBC

Despite the fact that this paper did not perform tests on SwarmBC, it is still possible to theoretically compare both SwarmAvg and SwarmBC, based on a comprehensive understanding of the inner workings of each.

One of the primary drawbacks of utilising SwarmAvg as opposed to SwarmBC is the potential occurrence of divergent training during SwarmAvg training. The detrimental effect of divergent training may persist throughout the entire training process, leading to significantly worse performance than anticipated. Several factors can lead to divergent training, with low settings for α and γ , coupled with a sparsely-connected network, being highly correlated with its occurrence during testing. SwarmBC does not suffer from this issue, as the blockchain ensures that all nodes agree on the same model. However, during testing, divergent training was a rare phenomenon, and its remedy was a simple parameter tuning process upon detection.

One drawback of utilising the SwarmAvg algorithm, as opposed to SwarmBC, is the absence of security features. The utilisation of blockchain technology enables the straightforward authentication of nodes utilising different algorithms that are frequently well-documented and integrated into the blockchain framework. On the other hand, SwarmAvg lacks any built-in authentication mechanism. As a result, to employ it securely, a separate authentication system would need to be implemented to prevent malicious parties from interfering with training.

SwarmAvg surpasses SwarmBC in terms of computational efficiency as nodes are not mandated to verify transactions with a blockchain, which involves a comparatively substantial overhead as opposed to simply averaging the models of a nodes neighbours. Consequently, SwarmBC necessitates higher processing power to operate for the same amount of training steps. This overhead could impede the speed of training on a resource-limited system, such as a swarm of drones, as a considerable amount of time is consumed in validating transactions, rather than performing training. Thus, SwarmAvg may be a more fitting choice than SwarmBC for swarms where processing power is a concern.

6.2 Evaluation of this Study

It must be acknowledged that this study presented certain limitations that necessitate careful consideration. The primary concerns are outlined as follows:

Small Simulated Number of Nodes: A significant concern regarding this study is the limited number of nodes simulated during testing. Several drawbacks of FL in comparison to SL, such as server bottlenecks, are not apparent when testing on a limited number of nodes.

Lack of Testing Overheads: In the course of testing the SwarmAvg algorithm, it was determined that the assessment of its performance should be based on training steps rather than time. Although it would have been preferable to

evaluate performance in relation to time, this was not feasible due to limitations in resources. The GPUs utilised in this study are primarily optimized for running a single GPU program effectively, whereas this research required multiple nodes to perform training simultaneously. As a result, it was noted during testing that training times varied significantly, rendering time-based measurement impractical due to high levels of noise.

Lack of Simulated Internet Lag: The algorithm proposed is designed for usage over the internet. However, the study solely evaluated its performance within a simulated environment, which lacked a time gap between the sending of information from one node to its reception by another. The deliberate omission of this time gap in the simulation aimed to increase the speed of simulations. However, this critical aspect of the algorithm’s functionality remains untested as a result.

Chapter 7

Project Management

7.1 Time Management

To assist with planning and organisation of this project, Gantt charts were used. These helped to visualise which tasks needed doing, and also if the project progress was ahead or behind schedule. The Gantt for the interim report can be found in appendix C.1 and the final report Gantt can be found at C.2.

On the Gantt charts, it was decided to not count weekends as working days. This was chosen as it represented the fact that the author had other work to do during the course of the project.

The presented charts represent the final iteration of planning. As described in the following section, changes occurred throughout the project and the Gantt charts were adapted accordingly.

7.2 Changing Plans

Initially, the project proposal included the implementation of both SL and a simulator. The intended purpose of the simulator was to generate camera feeds that simulated those of security and doorbell cameras from various locations in a neighbourhood. Subsequently, the SL model would be employed to learn the task of identifying car crashes. However, this project was deemed excessively ambitious, and as a result, a decision was taken to significantly streamline the project by discarding many of its components.

Plans were also changed on a smaller scale somewhat regularly, due to the experimental nature of the project. However, this was not an issue due to the inclusion of many buffer zones in the planned time allocation, such as *Bugfixing Time*. This meant that an overrun task could be finished during this time, so it would not effect the project moving forward.

7.3 Risk Assessment

Below is the risk assessment table created at the beginning of the project:

Risk	Severity	Likelihood	Score	Mitigation
Personal issues such as illness prevent work from being done	3	3	9	The code will be designed such that there are many checkpoints with acceptable results. This means that even if the project is not finished, it will still have an satisfactory outcome.
Failure of authors computer	4	2	8	Project is backed up on Github. Development environment is backed up on google drive using docker. This means the author has the possibility to run the code on the Zepler labs computers using an identical environment.
Algorithm does not work as expected	5	1	5	During research, the author read a large amount of literature on Distributed Federated Learning (DFL). It would be possible to design a DFL algorithm and perform tests on that instead, which is more likely to work.

Chapter 8

Conclusion and Future Work

8.1 Conclusion

This study has presented a novel SL algorithm, called SwarmAvg, which enables completely distributed and decentralised machine learning to take place. The algorithm is similar to FL in that it allows each node in the network to hold its own confidential dataset which never needs to be shared with any other nodes.

It has been found that SwarmAvg is less performant than FedAvg in all situations where FedAvg can be applied. However, in many of these scenarios, SwarmAvg reached a similar final accuracy to FedAvg, but training took more iterations to reach this. SwarmAvg showed most potential when performing training in a sparsely connected network of nodes, a situation where FedAvg cannot function, but SwarmAvg saw only minor reductions in performance when compared to a densely connected network. Unfortunately, it was shown that SwarmAvg performs significantly worse than FedAvg in situations where each node only has access to a small subset of the total classes. Despite the fact that this is an extreme case, unevenly distributed data is an issue for many algorithms in the real world, possibly indicating that SwarmAvg may need further improvements before deployment.

Notwithstanding these disadvantages, SwarmAvg could potentially be a compelling algorithm to employ in particular circumstances. For instance, in a very large swarm, FL may become impossible to use due to the bottleneck of sending all data to a single node, a problem that SwarmAvg does not face. SwarmAvg also may be a better choice if it is known that the network in which the algorithm will be deployed to is very unreliable, with nodes and connections dropping out.

Overall, the goals specified in Section 1.3 have been met, so the project can be considered a success:

- (A) A novel algorithm for use with machine learning was successfully designed. It functioned in a completely decentralised manner.
- (B) The algorithm was successfully implemented.
- (C) The performance was tested in densely connected networks, with both access to the full dataset and access to a subset of the dataset's classes. It was compared to FedAvg.

- (D) The algorithm was shown to perform in situations where FedAvg could not be used.

8.2 Future Work

Simulate More Nodes: In this paper, only 10 concurrent nodes were tested. However, if an internet-enabled back-end was to be implemented, it would be feasible to distribute numerous nodes among multiple training machines, thus enabling the simulation of larger swarms.

Study of Effects of SwarmAvg Parameters: For the present study, parameters were chosen heuristically to perform testing, potentially resulting in sub-optimal outcomes as compared to what could have been achievable. An in depth examinations into the effects of different parameters could be beneficial to future work.

Experimenting with a Hybrid FL-SL approach: In this study, FedAvg has been shown to be more effective in scenarios where the environment is tightly controlled when compared to SwarmAvg. However, a potential approach would multiple utilise small clusters of FedAvg, where the FedAvg servers are connected with SwarmAvg. This may inherit both the high performance of FedAvg, but with some of the fault tolerance of SwarmAvg.

Dynamic Environment: In all of the presented experiments, the simulated environments were static, meaning that the connections between nodes did not change throughout training. A dynamic environment may have significant advantages, as each node is exposed to more varied neighbours, which could lead to higher performance.

Improve Unevenly Distributed Class Performance: The scenarios in which SwarmAvg was tested where each node had access to only a subset of the classes left much to be desired. SwarmAvg performed far worse than FedAvg in these situations, however it remains plausible that with future research and modifications, that SwarmAvg may be able to handle these situations better.

Bibliography

- [1] M. Kop, “Machine learning & eu data sharing practices,” Stanford-Vienna Transatlantic Technology Law Forum, Transatlantic Antitrust . . . , 2020.
- [2] G. A. Kaissis, M. R. Makowski, D. Rückert, and R. F. Braren, “Secure, privacy-preserving and federated machine learning in medical imaging,” *Nature Machine Intelligence*, vol. 2, no. 6, pp. 305–311, 2020.
- [3] W. L. Shuya Ke, “Distributed multi-agent learning is more effectively than single-agent,”
- [4] C. Zhang, Y. Xie, H. Bai, B. Yu, W. Li, and Y. Gao, “A survey on federated learning,” *Knowledge-Based Systems*, vol. 216, p. 106775, 2021.
- [5] Q. Li, Z. Wen, Z. Wu, S. Hu, N. Wang, Y. Li, X. Liu, and B. He, “A survey on federated learning systems: vision, hype and reality for data privacy and protection,” *IEEE Transactions on Knowledge and Data Engineering*, 2021.
- [6] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas, “Communication-efficient learning of deep networks from decentralized data,” 2016.
- [7] V. Mothukuri, R. M. Parizi, S. Pouriyeh, Y. Huang, A. Dehghantanha, and G. Srivastava, “A survey on security and privacy of federated learning,” *Future Generation Computer Systems*, vol. 115, pp. 619–640, 2021.
- [8] T. Zhang, L. Gao, C. He, M. Zhang, B. Krishnamachari, and A. S. Avestimehr, “Federated learning for the internet of things: Applications, challenges, and opportunities,” *IEEE Internet of Things Magazine*, vol. 5, no. 1, pp. 24–29, 2022.
- [9] M. Xie, G. Long, T. Shen, T. Zhou, X. Wang, and J. Jiang, “Multi-center federated learning,” *CoRR*, vol. abs/2005.01026, 2020.
- [10] J. Guo, Q. Zhao, G. Li, Y. Chen, C. Lao, and L. Feng, “Decentralized federated learning with privacy-preserving for recommendation systems,” *Enterprise Information Systems*, vol. 0, no. 0, p. 2193163, 2023.
- [11] D. Yaga, P. Mell, N. Roby, and K. Scarfone, “Blockchain technology overview,” tech. rep., oct 2018.
- [12] D. Yang, C. Long, H. Xu, and S. Peng, “A review on scalability of blockchain,” in *Proceedings of the 2020 The 2nd International Conference on Blockchain Technology*, ICBCT’20, (New York, NY, USA), p. 1–6, Association for Computing Machinery, 2020.

- [13] W.-H. et al., “Swarm learning for decentralized and confidential clinical machine learning,” *Nature*, vol. 594, pp. 265–270, Jun 2021.
- [14] J. C. Varughese, R. Thenius, T. Schmickl, and F. Wotawa, “Quantification and analysis of the resilience of two swarm intelligent algorithms,” in *GCAI*, pp. 148–161, 2017.
- [15] J. Augustine, T. Kulkarni, and S. Sivasubramaniam, “Leader election in sparse dynamic networks with churn,” in *2015 IEEE International Parallel and Distributed Processing Symposium*, pp. 347–356, IEEE, 2015.
- [16] M. Dorigo, M. Birattari, and M. Brambilla, “Swarm robotics,” *Scholarpedia*, vol. 9, no. 1, p. 1463, 2014. revision #138643.
- [17] J.-H. Chen, M.-R. Chen, G.-Q. Zeng, and J.-S. Weng, “Bdfl: a byzantine-fault-tolerance decentralized federated learning method for autonomous vehicle,” *IEEE Transactions on Vehicular Technology*, vol. 70, no. 9, pp. 8639–8652, 2021.
- [18] H. Xiao, K. Rasul, and R. Vollgraf, “Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms,” 2017.

Appendix A

Network Generation Algorithm

```
def count_connections(n):
    if n == 1:
        return 0
    return count_connections(n-1)+n-1

def is_connection_in(c, cons):
    return c in cons or (c[1], c[0]) in cons

def random_pair(length):
    a = np.random.randint(0, length)
    b = a
    while b == a:
        b = np.random.randint(0, length)
    return a, b

def fully_connected_graph(nodes_names, density=0):
    # At 0 density the graph is connected minimally
    # (all nodes connect to all nodes)
    # At 1 density the graph is fully connected
    nodes = [[x] for x in nodes_names]
    # Make a random fully connected network
    connections = []
    while len(nodes) > 1:
        # Pick the two random groups
        # They cannot be the same group
        a, b = random_pair(len(nodes))
        na, nb = nodes[a], nodes[b]
        # Pick the member of each group to connect
        ma = na[np.random.randint(0, len(na))]
        mb = nb[np.random.randint(0, len(nb))]
        # Add a connection between them
        connections.append((ma, mb))

    # Make a new group that has all nodes
    # from the two groups that have been connected
```

```

    newgroup = []
    for n in na:
        newgroup.append(n)
    for n in nb:
        newgroup.append(n)
    nodes = [n for n in nodes
              if n not in [nodes[a], nodes[b]]]
    nodes.append(newgroup)
# Add connections until we reach the density
    minimum_connections = len(connections)
    maximum_connections = count_connections(len(nodes_names))
    target_connections =
        int((maximum_connections - minimum_connections)*density)
        +minimum_connections
    while len(connections) < target_connections:
        con = connections[0]
        while is_connection_in(con, connections):
            a, b = random_pair(len(nodes_names))
            con = (nodes_names[a], nodes_names[b])
        connections.append(con)
    return connections

```

Appendix B


Machine Learning Model

```
inp = Input((28,28))
out = Reshape((28,28,1))(inp)
out = Conv2D(16, (3,3), activation="relu")(out)
out = Conv2D(16, (3,3), activation="relu")(out)
out = Flatten()(out)
out = Dense(256, activation="relu")(out)
out = Dense(128, activation="relu")(out)
out = Dense(10, activation="sigmoid")(out)
model = Model(inputs=inp, outputs=out)
model.compile(
    optimizer="adam",
    loss=SparseCategoricalCrossentropy(),
    metrics=[SparseCategoricalAccuracy()]
)
```


Appendix C

Gantt Charts

C.1 Gantt - Interim



C.2 Gantt - Final



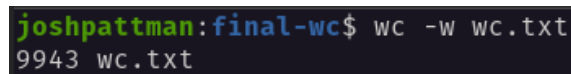
Appendix D

Verifying the Word Count

Below are the steps used to attain the verifiable word count.

1. Extract the pages starting from the first page of 'Introduction, Problems and Goals' down to the last page of 'Future Work'.
2. Copy all content from the pdf using a pdf viewer into a plaintext file. This ensures that pdf formatting symbols are not included. (Ctrl+A -> Ctrl+V).
3. Run the command `wc -w <name-of-text-file>`

When performing this method, the output of the terminal window is shown in Figure D.1:



```
joshpattman:final-wc$ wc -w wc.txt
9943 wc.txt
```

Figure D.1: The terminal window output for the word count command

This means that the word count for this document is verifiably 9943.

Appendix E

Data Archive Overview

Below is the file structure for the accompanying project archive.

```
project_archive
├── data: The location which data from running tests is stored
│   ├── test2.3.json: The data from test 2 iteration 3
│   └── ...
├── colored_log_formatter.py: (NOT MY CODE BUT NEEDED TO RUN) Helper functions for colouring the python logger
├── experiments.sh: File to re-run all experiments
├── fed_mnist.py: File to run an experiment with FL
├── fed: Contains files for FL algorithms
│   ├── backend.py: FL algorithm functions
├── swarm_mnist.py: File to run an experiment with SwarmAvg
├── fed: Contains files for SwarmAvg algorithms
│   ├── local_backend.py: System to simulate a network
│   └── swarm_dist.py: SwarmAvg algorithm
├── flatten_model.py: Helper function to flatten/unflatten a Keras model
├── graphs.py: Code to generate sparse graphs
├── just_model_mnist.py: Test file to train the model with single node
├── model.py: File to generate the machine learning model
└── visualise.ipynb: File to generate graphs from data in ./data/
```

Appendix F

Original Project Brief

Privacy Focussed Self Improving Disaster Detection Using A Distributed Network

Josh Pattman - Supervised by Mohammad Soorati

October 25, 2022

1 Problem

Disasters, such as house fires and car crashes, are an unfortunate yet frequent part of our lives in the modern world. Today, we still rely heavily upon passers by to call emergency services, and dispatch help to the situation. However, this is not only slow, but it can also be difficult to understand a panicked citizen who could give the emergency services incorrect information in their fluster. In a disaster situation, every second counts, so any improvement to this system could save lives.

2 Proposal

From this point on an internet connected device with a camera and known position and orientation will be referred to as an agent.


2.1 Goals

Today, there are a multitude of devices that are capable of acting as agents. From smartphones to doorbell cameras to consumer drones, most areas of a modern city are under constant surveillance. However, due to privacy protection, it is inviable to stream this video to any location other than the device. The goals of this project are to: Report disasters as fast as possible, protect the privacy of any user who's device is being used as an agent, improve the systems response speed over time, and finally have exceptional fault tolerance.

2.2 Scope

This project will include development of the machine learning and fusing algorithm. It will provide access to these through simple command line tools and python notebooks. There will be no user interface created that could be used in real life, as this is out of scope of the project. Also, for a large portion of the experimentation phase, a simulation will be used to generate camera feeds around an area. At later stages, real life images may be used, however these will be far harder to come by, so may not be possible. The project will also entail running the algorithms on one or multiple computers on the same network, however the code will not be implemented onto actual devices such as doorbells or phones.

2.3 Proposed solution



2.4 Phases

1. **Phase 1:** A photo-realistic simulation will be created. It will stream in-simulation camera feeds to the various agents.
2. **Phase 2:** A single agent will be trained on video from the simulation to predict locations of accidents in 3D space.
3. **Phase 3:** Various video feeds from the simulation will be converted to 3D disaster coordinates, using the multiple copies of the trained agent from *phase 2*. This data will be fed into the fusing algorithm to create a 3D map of disasters.
4. **Phase 3:** An automatic feedback system will be created where the network of agents predict the position of disasters, and the data is then sent back to the simulation. The simulation then confirms or denies the existence of each disaster, which is forwarded to each agent on the network. Each agent then learns from this.
5. **Phase 4:** Phase 3 will be repeated with real world data, however the agents from Phase 3 will be used as a starting point.